

# Data Wrangling of theses data in France

Amr MOHAMED

CY Tech/ CY Cergy Paris University  
mohamedamr@cy-tech.fr

## Abstract

In this paper, we are going to go through an analysis that was done over a real-life dataset of Ph.D. theses that were defended in France between the years 1971 and 2020, and we are going through all the steps of a data wrangling project starting with Web scraping the data from the internet, Data preprocessing, Data Cleaning, Spotting abnormalities and outliers in the data, and finally, we will draw assumptions using Data visualization and some statistical methods.

**Keywords:** Theses, Data-Wrangling, R

## 1 Introduction

During the past decade, Data Science has become a very important field in all industries as it shows us the real insights behind the different sources of data. In this paper, we will be drawing insights about Ph.D. theses that were defended between 1971 and 2020 in France through going through all the variables, identifying the missingness and abnormalities in the data, moreover, we will go deeply to analyze the trends of the numerical variables over time like how the number of theses changed over the past half of a century, and when was it high/low, in addition, we are going to see how we can verify our analysis from the data by Web-Scraping.

## 2 Data Source

The data that was used in this study was from theses.fr<sup>1</sup> website which was launched by France in the early 2000s. The dataset consists of 447644 rows and 18 columns. Each row in the dataset represents a defended thesis and the columns define all the information about the thesis author, supervisor(s), university (institute), and dates of defense and inscription.

## 3 Data Scraping

Firstly, we started by scraping the data from theses.fr website using python library 'Beautiful Soup' for web scraping. The structure of the data is 18 columns and 447644 rows. The scraped columns are represented in Table.1 below.

Column	Description
Auteur	String Author name
Identifiant.auteur	Integer author ID
Titre	String title of the thesis
Directeur.de.these	String name(s) of the supervisor(s)
Directeur.de.these.nom.prenom	String last.name first.name of the supervisor(s)
Identifiant.directeur	Integer ID of the supervisor(s)
Etablissement.de.soutenance	String name of the Institute(University)
Identifiant.etablissement	Integer ID of the institute
Discipline	String Domain name
Statut	String thesis status
Date.de.premiere.inscription.en.doctorat	Date of first inscription in the doctorate
Date.de.soutenance	Date, Date of the defense of the thesis
Year	Date, year of the defense of the thesis
Langue.de.la.these	String, Language of the thesis
Identifiant.de.la.these	Integer, ID of the thesis
Accessible.en.ligne	Boolean yes/no of the accessibility online
Publication.dans .theses.fr	Date, Date of the publication on theses.fr
Mise.a.jour.dans .theses.fr	Date, Date of the last update on theses.fr

Table 1: Columns of the dataset description

<sup>1</sup><https://theses.fr/>

## 4 Dealing with Missing Data

Afterward, we started inspecting the missingness in the dataset. Firstly, the function `vis_miss()` from `naniar` library to visualize the missing data in each column and get its percentages.




Figure 1: Missing data percentage per column

Column	missing	%
Date.de.premiere.inscription.en.doctorat	383668.00	86.00
Identifiant.auteur	129989.00	29.00
Langue.de.la.these	63765.00	14.00
Date.de.soutenance	56746.00	13.00
Year	56746.00	13.00
Identifiant.directeur	49172.00	11.00
Identifiant.etablissement	17085.00	4.00
Mise.a.jour.dans.theses.fr	177.00	0.00
Directeur.de.these	15.00	0.00
Directeur.de.these.nom.prenom.	15.00	0.00
Titre	9.00	0.00
Discipline	5.00	0.00
Etablissement.de.soutenance	4.00	0.00
Auteur	0.00	0.00
Statut	0.00	0.00
Identifiant.de.la.these	0.00	0.00
Accessible.en.ligne	0.00	0.00
Publication.dans.theses.fr	0.00	0.00

Table 2: Missing data count and percentage per column

From Figure.1 and Table.2, we can see that the column 'Date.de.premiere.inscription.en.doctorat' has the biggest proportion of missing values with a percentage of 86% followed by 29% of missing data in the 'Identifiant.auteur' column, and then the missing values decrease for 14-13% for the columns 'Langue.de.la.these', 'Date.de.soutenance', 'Year', 'Identifiant.directeur', and

4% for the column 'Identifiant.etablissement' and the rest of the columns approximately have 0% of missing values.

We can notice from the graph along with the table above that the columns 'Date.de.soutenance' and 'Date.de.premiere.inscription.en.doctorat' are highly dependent, where they complement each other in the sense that the rows which have missing values in 'Date.de.soutenance' aren't missing in 'Date.de.premiere.inscription.en.doctorat' and vice versa, and this can be explained by the fact that the thesis is officially considered as a Ph.D. thesis by the french education system on the day on which it was defended.

Moreover, a new variable was created using `rnorm()` function to simulate the number of pages per thesis, with a mean number of pages of 200 pages and a standard deviation of 50 pages, where 20% of data in the created column is missing. To treat this missingness in the number of pages column, we used imputation with the mean number of pages as the data is centered around the mean, and the missingness is limited, so that the imputation will not affect the mean of the data, and won't cause many effects on the standard deviation.

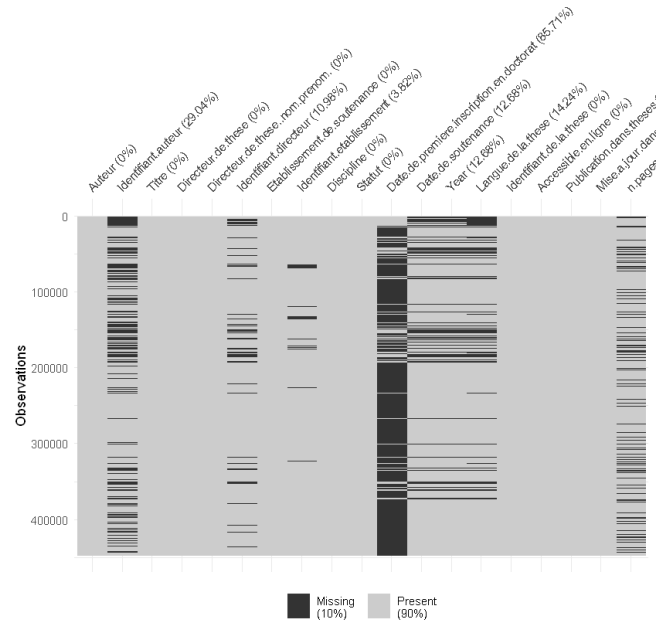


Figure 2: Missing data per column percentages before the number of pages imputation




Figure 3: Missing data per column percentages after n.pages imputation

As we can see from Figure 2 and Figure 3, the number of pages imputation was made successfully and reduced the total missingness of data by 1.1%.

## 5 Common Issues

In this section, we are going to go through the common issues that were discovered during the data wrangling process and what can be the causes of such issues.

Firstly, during the processing of the 'Date.de.soutenance' column of the dataset, it was found that 71.8% of the theses in the dataset were recorded to be defended on the first of January over the years. These insights were found using the function `filter()` and `group_by()` from the library `dyplyr` in R, where all the dates were split into 3 columns of the day, month, and year then grouped and counted the theses that were defended on the first day of the first month of each year. From a primary point of view, it is clear that the theses which were found to be defended on the first of January were recorded on a default date which is the first of January and that is why we can see the majority of the theses are defended on that date, while in fact, this day is a public holiday in France, and no theses can be defended on that day.




Figure 4: Percentage of theses on the first of January over years

As we can see in Figure 4, the percentage on the theses recorded to be defended on the first of January over the years decreased over time from being approximately equal to 100% in 1971 till it became approximately equal to 0% in 2020.

Moreover, in the 'Auteur' column of authors' names, we have found multiple theses that were defended by the same name, so we had to investigate the reason why those names appeared multiple times in the dataset.

We picked the name 'Cécile Martin' from the dataset which had 7 theses defended in total. We applied `group.by()` and `filter()` methods along with `count` on the author name and the author ID.

Author	ID	Frequency
Cecile Martin	179423568	1
Cecile Martin	182118703	1
Cecile Martin	203208145	1
Cecile Martin	81323557	4

Table 3: Theses associated with author name Cécile Martin

From Table.3 above, it was found that the 7 theses associated with the name Cécile Martin were defended by 4 different authors, and Cécile Martins with ID (81323557) has defended 4 theses and it was verified from the dataset that the 4 theses are all in fields associated with Biology. Moreover, it was verified through theses.fr that Cécile Martins with ID (81323557) defended the 4 theses <sup>2</sup>.

For further verification, we built a scraper function from scratch to access any author’s profile on theses.fr website, and scrap the number of theses defended by the author, and so we can verify whether the author has defended the number of theses given in the dataset or it is a mistake. In Cécile Martin’s case, the scraper returned that Cécile Martin has defended 4 theses, so the result we found through the dataset is correct.

Afterward, we proceeded to check for any issues in the 'Identifiant.directeur' column which represents the ID of the supervisor(s) who supervised the theses. Firstly, we found 3 main issues in the supervisor id columns which are:

- 35163 supervisors' IDs end with the character 'X'.
- 59108 supervisors' IDs have commas in the IDs.
- 4724 supervisors' IDs consist of 1 or 2 digit numbers.

<sup>2</sup><https://theses.fr/081323557>

After further investigation, it was found that the supervisors' IDs that end with the character 'X' can be treated by removing the character 'X' from the end of the IDs. Then we investigated the IDs containing commas in them, and it was found that those IDs are incorrectly recorded as a reason that those theses were supervised by multiple supervisors.

In addition, we observed a severe recession in the number of theses defended in the years 2019 and 2020.




Figure 5: Number of theses defended over the years

As we can see in Figure 5, there is a severe recession in the number of theses defended per year. From a preliminary point of view, we suppose to be caused by COVID-19 which forced most of the academic institutions to be closed for approximately 1 year and a half.

## 6 Outliers

After spotting the common issues in the main columns of the dataset, we went further to check whether the data has outliers or not and if it has if these outliers recorded y mistake.

We started by investigating the supervisors by checking the number of theses that they have mentored. We used `group_by()` method along with `filter()`, `count()`, and `arrange()` to get the most frequent supervisors in the dataset.

	Supervisor	Supervisor ID	Frequency
1	Jean-Michel Scherrmann	59375140	208
2	Francois-Paul Blanc	26730774	205
3	Pierre Brunel	26756625	193
4	Philippe Delebecque	29561248	178
5	Guy Pujolle	27084868	177
6	Michel Bertucat	98531891	173
7	Bernard Teyssie	27158578	146
8	Bruno Foucart	26870177	132
9	Henry de Lumley	26997894	132
10	Jean-Claude Chaumeil	58552499	131

Table 4: Supervisors with the highest numbers of these supervisions

From Table 4. above, we can see that the supervisors above have mentored a very large number of theses. To verify that the supervisors in the table have mentored these numbers of theses, we used `group_by()` with the supervisors' names and IDs to make sure that they

weren't recorded incorrectly. Unexpectedly, those numbers were correct, and the supervisors mentored these large numbers of theses.

To verify the number of theses supervised by each supervisor from Table 4, we built a scraper function, similar to the one which was built for Authors' number of theses defended verification, for supervisors, to be able to verify the number of theses mentored by each supervisor. By running the scraper we got an accuracy of 72% of the outliers to have correct values, and the rest are even mistakes or have supervised more theses after the collection of the data.

To verify our findings statistically, we used the `quantile()` function to get the extreme values below and above which the number of supervised theses by the supervisor is considered as an outlier. For the 97.5 and 2.5 percentiles, we got that the values of 13 and 1 respectively, which means that any number of these supervised less than 1 or greater than 13 is an outlier. Moreover, we calculated the z-score for the number of supervised theses with a 95% confidence interval, we got approximately the same result where it showed that if the number of theses is greater than or equal to 12 or less than 1, it is considered as an outlier.

Moreover, we used the same techniques used earlier in identifying the outliers in supervisors on the column of authors. After grouping the data by the authors' names and IDs, we used the quantiles method to detect the range in which below or above it the author's number of theses defended is considered as an outlier. Surprisingly, specifying the 2.5 and 97.5 percentiles, it was found that any author who has several theses defended different than 1 thesis, is considered as having an outlier number of theses defended.

And for further verification, we used the scraper built earlier to verify the number of theses defended by Cecile Martin to verify the outliers in the authors. When the scraper was run on the first 15 authors with the highest number of these defenses, it gave an accuracy of 76

To verify the findings of the number of theses defended by the authors, we calculated the z-score to find that with more than 99% confidence interval, any author who defended different than 1 thesis, is considered as having an outlier number of theses defended.

## 7 Preliminary Results

Afterward, we moved forward to process the languages. We started by recoding the languages data using the `dplyr` library into four categories which are: French, English, Bilingual, and Other, where a thesis is considered bilingual if it is written in English and French, and the Other category is for any other condition which is not in the first 3. We started by recoding the languages using the `mutate()` function from `dyplr` package in R, then we used `filter()`, `group_by()`, and `count()` to get the count of theses written in each language category.

From table 5 above we can see the frequency of each language category in the dataset, and it's clear through

	Language category	Frequency
1	Bilingual	10576
2	English	30942
3	French	334404
4	Other	71722

Table 5: Language categories frequencies

the frequencies that French is the most used language in writing the theses in France, which is very reasonable since these theses were written in France.




Figure 6: Number of theses per language category

To have a better look at the frequencies of the language categories in the dataset, we plotted the bar graph above, and it enhances our observation of the French language is the most used in writing the theses in France.

In addition, we were interested in investigating how the choice of the language of the manuscript evolved over the past decades. So, we used the library dplyr to filter, group, and count the theses that were defended on the same period, so that we can see the trend.




Figure 7: Number of theses per language category over the years

As we can see in Figure 7 surprisingly, the trend of the English category started to increase significantly over the past two decades.

In addition, we were interested in knowing in which month do the authors tend to defend their theses. As we did previously, we used dplyr to group, filter, and count the data.




Figure 8: Average percentages of theses defended per month

As we can see in Figure 8, approximately 27% of the authors tend to defend their theses in December, while approximately 0% tend to defend their theses in August.

Finally, we were interested in knowing the gender of authors that defended their theses over the past decades. To get this information, we processed the author names column in python and ran it through a library that detects the gender from the name. Then imported the generated data in R, and performed the previous steps of filtering, grouping, and counting.




Figure 9: Number of theses defended over the years by gender of authors

From Fig 9, we can see the trend where male and female authors have almost the same distribution of the number of theses over the years, but with fewer numbers for the males with approximately 25% less.

Moreover, we were interested to see the same trend of the number of theses supervised by gender over the year for the supervisors, so repeated the same steps which were made to produce the Previous figure.




Figure 10: Number of theses supervised over the years by gender of supervisors

As seen before in Figure 10, In Figure 11 we can see the same trend over the years for female and male supervisors for the supervised theses, but this time with almost an equal number of theses supervised by both genders.

## 8 Appendix




Figure 11: Number of theses per language category over the years

## References

Theses website: <https://theses.fr/>

Plotly area Plots: <https://plotly.com/r/filled-area-plots/>

STHD ggplot area plot: <http://www.sthda.com/english/wiki/ggplot2-area-plot-quick-start-guide-r-software-and-data-visualization>

Visualizing missing data with Naniar : <https://cran.r-project.org/web/packages/naniar/vignettes/getting-started-w-naniar.html>